

Large Language Models (LLMs) are designed to generate human-like text based on the data they were trained on. However, LLMs can sometimes produce hallucinated responses, meaning the generated information is not grounded in the provided source material.

Hallucination detection systems aim to evaluate whether an LLM-generated answer is supported by evidence. Such systems commonly use document retrieval, semantic similarity, and natural language inference (NLI) to determine if an answer is factual, partially supported, or contradicted by source documents.

Trust scoring frameworks assign numerical scores to model outputs, helping users assess reliability. These scores are especially important in domains such as healthcare, law, and education.